



# Machine Learning 10-601

Tom M. Mitchell  
Machine Learning Department  
Carnegie Mellon University

December 4, 2017

Today:

Review for final exam

Recommended reading:

- Mitchell: “Key Ideas in ML”

# Loss Functions

## AdaBoost Algorithm

**Given:**  $(x_1, y_1), \dots, (x_m, y_m)$  where  $x_i \in X, y_i \in Y = \{-1, +1\}$

**Initialize**  $D_1(i) = 1/m$ .

**For**  $t = 1, \dots, T$ :

- Train weak learner using distribution  $D_t$ .
- Get weak hypothesis  $h_t : X \rightarrow \{-1, +1\}$  with error

$$\epsilon_t = \Pr_{i \sim D_t} [h_t(x_i) \neq y_i].$$

- Choose  $\alpha_t = \frac{1}{2} \ln \left( \frac{1 - \epsilon_t}{\epsilon_t} \right)$ .
- Update:

$$\begin{aligned} D_{t+1}(i) &= \frac{D_t(i)}{Z_t} \times \begin{cases} e^{-\alpha_t} & \text{if } h_t(x_i) = y_i \\ e^{\alpha_t} & \text{if } h_t(x_i) \neq y_i \end{cases} \\ &= \frac{D_t(i) \exp(-\alpha_t y_i h_t(x_i))}{Z_t} \end{aligned}$$


where  $Z_t$  is a normalization factor (chosen so that  $D_{t+1}$  will be a distribution).

**Output** the final hypothesis:

$$H(x) = \text{sign} \left( \sum_{t=1}^T \alpha_t h_t(x) \right).$$

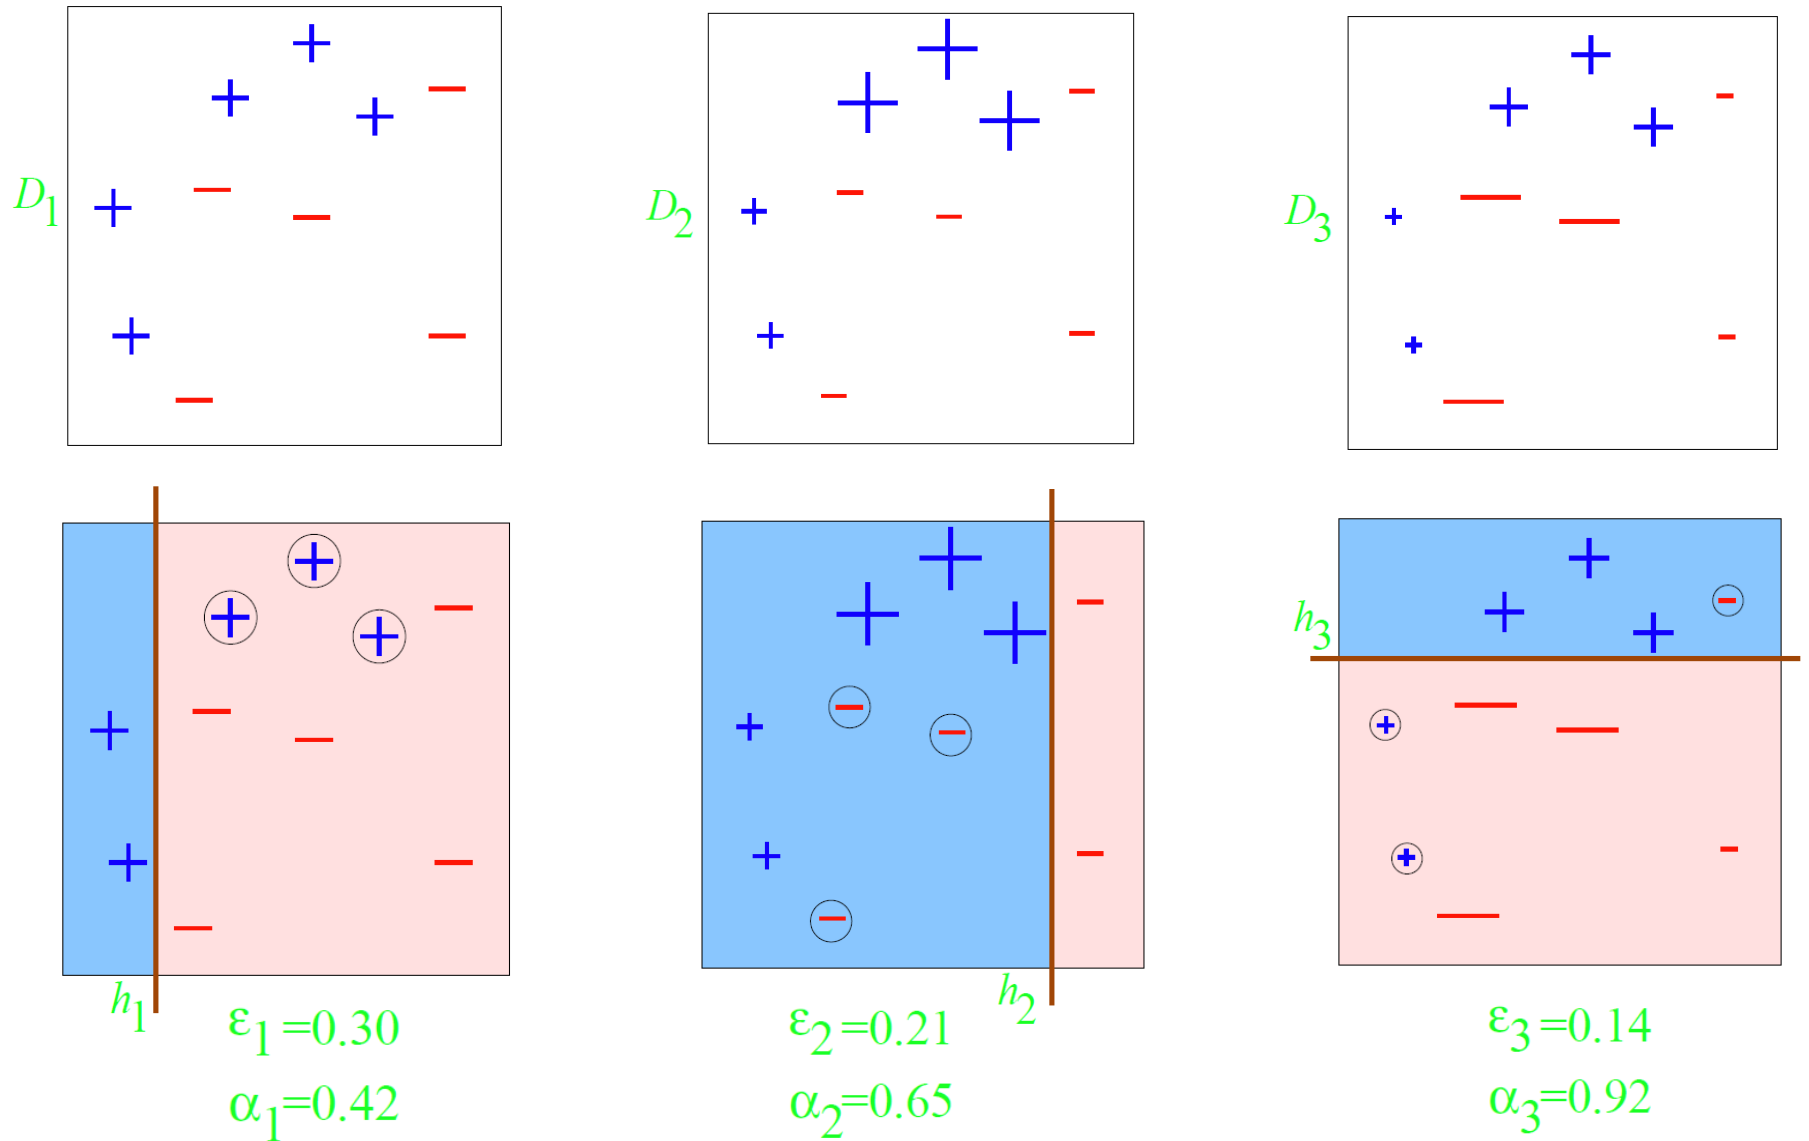
# AdaBoost: A toy example

Weak classifiers: vertical or horizontal half-planes (a.k.a. decision stumps)




[Rob Schapire]

# AdaBoost: A toy example



# AdaBoost: A toy example



# Theoretical Result 1: Training Error

We can minimize this bound by choosing  $\alpha_t$  on each iteration to minimize  $Z_t$

$$Z_t = \sum_{i=1}^m D_t(i) \exp(-\alpha_t y_i h_t(x_i))$$

For boolean target function, this is accomplished by [Freund & Schapire '97]:

$$\alpha_t = \frac{1}{2} \ln \left( \frac{1 - \epsilon_t}{\epsilon_t} \right)$$

Where:

$$\epsilon_t = \sum_{i=1}^m D_t(i) \delta(h_t(x_i) \neq y_i)$$

# Logistic Regression

Logistic regression assumes:

$$P(Y = 1|X = \mathbf{x}) = \frac{1}{1 + \exp(-(\mathbf{w}\mathbf{x} + b))}$$

Equivalently if  $y \in \{-1, +1\}$  :

$$P(Y = y|X = \mathbf{x}) = \frac{1}{1 + \exp(-y(\mathbf{w}\mathbf{x} + b))}$$

And trains to minimize log loss:

$$loss = \sum_{j=1}^m \ln(1 + \exp(-y_j(\mathbf{w}\mathbf{x}_j + b)))$$





# Logistic regression and Boosting

## Logistic regression:

- Minimize loss fn

$$\sum_{i=1}^m \ln(1 + \exp(-y_i f(x_i)))$$

- Define

$$f(x) = \sum_k w_k x_k + b$$

where  $x_j$  predefined

## Boosting:

- Minimize loss fn

$$\sum_{i=1}^m \exp(-y_i f(x_i))$$

- Define

$$f(x) = \sum_t \alpha_t h_t(x)$$

where  $h_t(x_i)$  defined  
dynamically to fit data

- Weights  $\alpha_j$  learned  
incrementally

# Boosting and Logistic Regression

Logistic regression equivalent to minimizing log loss


$$\sum_{i=1}^m \ln(1 + \exp(-y_i f(x_i)))$$

Boosting minimizes similar loss function!!

$$\frac{1}{m} \sum_i \exp(-y_i f(x_i)) = \prod_t Z_t$$

**Both smooth approximations of 0/1 loss!**

# Slack variables – Hinge loss



Complexity penalization


$$\xi_j = \text{loss}(f(x_j), y_j)$$



$$f(x_j) = \text{sgn}(\mathbf{w} \cdot \mathbf{x}_j + b)$$

$$\begin{aligned} \min_{\mathbf{w}, b} \quad & \mathbf{w}^T \mathbf{w} + C \sum_j \xi_j \\ \text{s.t.} \quad & (\mathbf{w}^T \mathbf{x}_j + b) y_j \geq 1 - \xi_j \quad \forall j \\ & \xi_j \geq 0 \quad \forall j \end{aligned}$$

$$\xi_j = (1 - (\mathbf{w} \cdot \mathbf{x}_j + b)y_j)_+ \quad \leftarrow \text{Hinge loss}$$




# SVM vs. Logistic Regression

## SVM : **Hinge loss**

$$\text{loss}(f(x_j), y_j) = (1 - (\mathbf{w} \cdot x_j + b)y_j)_+$$

## Logistic Regression : **Log loss** (negative log conditional likelihood)

$$\text{loss}(f(x_j), y_j) = -\log P(y_j \mid x_j, \mathbf{w}, b) = \log(1 + e^{-(\mathbf{w} \cdot x_j + b)y_j})$$




Boosting :  $loss(\mathbf{x}_j, y_j) = \exp(-(y_j \sum_t \alpha_t h_t(\mathbf{x}_j)))$

SVM : **Hinge loss**

$$loss(\mathbf{x}_j, y_j) = (1 - (\mathbf{w} \cdot x_j + b)y_j)_+$$

Logistic Regression : **Log loss** (negative log conditional likelihood)

$$loss(\mathbf{x}_j, y_j) = -\log P(y_j | x_j, \mathbf{w}, b) = \log(1 + e^{-(\mathbf{w} \cdot x_j + b)y_j})$$



# Loss Functions

- MLE, MAP
- log loss = negative log likelihood
  - minimize log loss = minimize misassigned probability mass
- regression: sum of squared errors
- regression: sum of squared errors plus square of weights
- SVM Hinge loss (maximize margin with slack variables)
- AdaBoost
- Bayes nets
- Logistic regression

# What You Should Know

---

- Sample complexity varies with the learning setting
  - Learner actively queries trainer
  - Examples arrive at random
  - ...
- Within the PAC learning setting, we can bound the probability that learner will output hypothesis with given error
  - For ANY consistent learner (case where  $c \in H$ )
  - For ANY “best fit” hypothesis (agnostic learning, where perhaps  $c$  not in  $H$ )
- VC dimension as measure of complexity of  $H$
- Mistake bounds
- Conference on Learning Theory: <http://www.learningtheory.org>
- ML Department course on Machine Learning Theory



# Kernels : Key Points

- Many learning tasks are framed as optimization problems
- Primal and Dual formulations of optimization problems
- Dual version framed in terms of dot products between  $x$ 's
- Kernel functions  $k(x,y)$  allow calculating dot products  $\langle \Phi(x), \Phi(y) \rangle$  without bothering to project  $x$  into  $\Phi(x)$
- Leads to major efficiencies, and ability to use very high dimensional (virtual) feature spaces

# What you should know

Primal and Dual optimization problems

Kernel functions

Support Vector Machines

- Maximizing margin
- Kernel SVM's
- Noise, slack variables and hinge loss
- Relationship between SVMs and logistic regression
  - 0/1 loss
  - Hinge loss
  - Log loss

Theory shows overfitting, mistakes depends on margin size

# What you should know

1. Using unlabeled data to reweight labeled examples gives better approximation to true error
  - If we assume examples drawn from fixed  $P(X)$
2. Unlabeled can help EM learn Bayes nets for  $P(X,Y)$ , and thus  $P(Y|X)$ 
  - If we assume the Bayes net structure reflects cond. independencies

## 2.5. Transductive SVM's

- If we assume maximizing margin captures relationship between  $P(X)$  and  $f: X \rightarrow Y$
3. Jointly train multiple classifiers, coupled by consistency constraints that can be evaluated using unlabeled data
    - optimize both the fit to labeled examples, and satisfaction of the consistency constraints

# What You Should Know

---

- Ensemble methods
- Weighted Majority
  - Learns weights for a given pool of hypotheses
  - Mistake bound relative to best hypothesis in the pool
  - ...
- Boosting
  - Learns weights and hypotheses
  - Theory: training error, true error, correspondence to Log. Regression
  - Practice: Boosted decision trees (and stumps) very popular!
- Many variants of ensemble methods
  - Resample training data to generate variety
  - Randomize learning algorithm to generate variety
  - Active learning – choose examples where vote is closest to tie

# Representation Learning: What you should know

- Unsupervised dimension reduction using all features
  - Principle Components Analysis
    - Minimize reconstruction error
  - Singular Value Decomposition
    - Efficient PCA
  - Independent components analysis
  - Canonical correlation analysis
  - Probabilistic models with latent variables
- Supervised dimension reduction
  - Hidden layers of Neural Networks
    - Most flexible, local minima issues
- LOTS of ways of combining discovery of latent features with classification tasks

# Bias, Variance, Unavoidable Error


Bias in stats bias  $\hat{\theta}$  is  $E[\hat{\theta}] - \theta$

Bias in ML involves estimating  $\hat{f}_n$  <sup>take over  $P(x)$</sup>

<sup>incomplete H</sup>  
alg choosing h in H has biased pref (eg. short series)

Variance — due to finite samples of training data  
statistical anomalies, unrepresentative data

Unavoidable error for nondeterministic fns.



# Transforming Inputs, Then Classifying

- Neural net, kernel SVM, AdaBoost, WeightedMajority



# Perspectives on ML

- Optimization
- Probabilistic modeling
- Parameterized programs
- Evolution

optimal policy  $\pi^*$

$$Q(s, a) = r(s, a) + \gamma \max_{a'} Q(s', a')$$

$\pi^*: S \rightarrow \text{action}$

$Q(s, a)$  is the value of the state-action pair  $(s, a)$  under the optimal policy  $\pi^*$ . The equation shows that the value of a state-action pair is the immediate reward  $r(s, a)$  plus the discounted value of the next state  $s'$ , where the discount factor  $\gamma$  is applied to the next state's value.

$$f: X \rightarrow Y$$

$$f: \text{State} \rightarrow \text{action}$$

$$V^\pi: \text{State} \rightarrow \mathbb{R}$$

sum of discounted future rewards using  $\pi$  from this state

$$Q^\pi: \text{state} \times \text{action} \rightarrow \mathbb{R}$$



# Key Results

- No free lunch
- Sources of error: bias, variance, unavoidable error
- Overfitting
- Bayes nets
- Deep neural nets
- PAC learning theory
- Semi-supervised learning
- Learning ensembles of functions
- Representation learning
- Kernel methods
- Reinforcement learning

