



Machine Learning 10-601

Tom M. Mitchell
Machine Learning Department
Carnegie Mellon University

October 16, 2017


Today:

- Graphical models
- Inference
 - Markov Blanket
 - Gibbs sampling
- Learning
 - fully observed: MLE, MAP
 - partly observed data: EM

Readings:

- Bishop chapter 8, 9-9.2

Bayesian Networks Definition



A Bayes network represents the joint probability distribution over a collection of random variables

A Bayes network is a directed acyclic graph and a set of conditional probability distributions (CPD's)


- Each node denotes a random variable
- Edges denote dependencies
- For each node X_i its CPD defines $P(X_i | Pa(X_i))$
- The joint distribution over all variables is defined to be

$$P(X_1 \dots X_n) = \prod_i P(X_i | Pa(X_i))$$

$Pa(X)$ = immediate parents of X in the graph

Markov Blanket


The Markov blanket of a node x_i comprises the set of parents, children and co-parents of the node. It has the property that the conditional distribution of x_i , conditioned on all the remaining variables in the graph, is dependent only on the variables in the Markov blanket.



from [Bishop, 8.2]

Why Markov Blanket is Useful for Inference


$$P(h|a, c, j, g) = ?$$



let's use shorthand $P(a)$ to represent $P(A=a)$

Why Markov Blanket is Useful for Inference

$$P(h|a, c, j, g) = P(h|c, j)$$




$$\begin{aligned} P(h|c, j) &= \frac{P(h, c, j)}{P(c, j)} = \frac{P(c)P(h|c)P(j|h)}{P(c)P(h|c)P(j|h) + P(c)P(\neg h|c)P(j|\neg h)} \\ &= \frac{P(h|c)P(j|h)}{P(h|c)P(j|h) + P(\neg h|c)P(j|\neg h)} \end{aligned}$$

let's use shorthand $P(a)$ to represent $P(A=a)$

Generating a sample from joint distribution: easy

How can we generate random samples drawn according to $P(F,A,S,H,N)$?



To generate a random sample for roots of network (F or A):

1. let $\theta = P(F=1)$ # look up from CPD
2. r = random number drawn uniformly between 0 and 1
3. if $r < \theta$ then output $F=1$, else $F=0$

To generate a random sample for S , given F,A :

1. let $\theta = P(S=1|F=f,A=a)$ # use f, a from above step
2. r = random number drawn uniformly between 0 and 1
3. if $r < \theta$ then output $S=1$, else $S=0$

Continue, generating $H|S=s$ and $N|S=s$

Generating a sample from joint distribution: easy


Note we can estimate anything!

e.g., $P(N=n)$ by generating many samples from joint distribution, then count the fraction of samples for which $N=n$

Similarly, for anything else we care about

$$P(F=1|H=1, N=0)$$

→ weak but general method for estimating any probability term...



Gibbs Sampling:

Goal: Directly sample conditional distributions

$$P(X_1, \dots, X_n \mid X_{n+1}, \dots, X_m)$$

Approach:

- start with arbitrary initial values for unobserved $X_1^{(0)}, \dots, X_n^{(0)}$ (and the observed X_{n+1}, \dots, X_m)

- iterate for $s=0$ to a big number:

$$X_1^{s+1} \sim P(X_1 \mid X_2^s, X_3^s \dots X_n^s, X_{n+1}, \dots, X_m)$$


$$X_2^{s+1} \sim P(X_2 \mid X_1^{s+1}, X_3^s \dots X_n^s, X_{n+1}, \dots, X_m)$$

...

$$X_n^{s+1} \sim P(X_n \mid X_1^{s+1}, X_2^{s+1}, \dots, X_{n-1}^{s+1}, X_{n+1}, \dots, X_m)$$

Eventually (after burn-in), the collection of samples will constitute a sample of the true $P(X_1, \dots, X_n \mid X_{n+1}, \dots, X_m)$

* but often use every 100th sample, since iters not independent



Gibbs Sampling:

Approach:

- start with arbitrary initial values for $X_1^{(0)}, \dots, X_n^{(0)}$ (and observed X_{n+1}, \dots, X_m)
- iterate for $s=0$ to a big number:



$$X_1^{s+1} \sim P(X_1 | X_2^s, X_3^s \dots X_n^s, X_{n+1}, \dots, X_m)$$

$$X_2^{s+1} \sim P(X_2 | X_1^{s+1}, X_3^s \dots X_n^s, X_{n+1}, \dots, X_m)$$

...

$$X_n^{s+1} \sim P(X_n | X_1^{s+1}, X_2^{s+1}, \dots, X_{n-1}^{s+1}, X_{n+1}, \dots, X_m)$$

Only need Markov Blanket at each step!



Gibbs is special case of Markov Chain Monte Carlo method

Inference in Bayes Nets

- In general, intractable (NP-complete)
- For certain cases, tractable
 - Assigning probability to fully observed set of variables
 - Or if just one variable unobserved
 - Or for singly connected graphs (ie., no undirected loops)
 - Variable elimination
- Can often use Monte Carlo methods
 - Generate many samples, then count up the results
 - Gibbs sampling (example of Markov Chain Monte Carlo)
- Many other approaches
 - Variational methods for tractable approximate solutions
 - Junction tree, Belief propagation, ...

see Graphical Models course 10-708

What you should know: Inference in Bayes Nets

- In general, intractable (NP-complete)
- Probability for a given joint assignment
- Probability for one unobserved variable given all the others
- Conditional independence / D-separation
- Markov blanket
- How to use Markov blanket to simplify inference
- How to generate samples from joint distribution
 - and how to use samples to estimate anything
- Gibbs sampling

Learning Bayes Nets from Data

Learning of Bayes Nets

- Four categories of learning problems
 - Graph structure may be known/unknown
 - Variable values may be fully observed / partly unobserved
- Easy case: learn parameters when graph structure is *known*, and training data is *fully observed*
- Interesting case: graph *known*, data *partly observed*
- Gruesome case: graph structure *unknown*, data *partly unobserved*

Learning CPTs from Fully Observed Data

- Example: Consider learning the parameter


$$\theta_{s|ij} \equiv P(S = 1 | F = i, A = j)$$

- MLE (Max Likelihood Estimate) is

$$\theta_{s|ij} = \frac{\sum_{k=1}^K \delta(f_k = i, a_k = j, s_k = 1)}{\sum_{k=1}^K \delta(f_k = i, a_k = j)}$$

kth training example

$\delta(X) = 1$ if X is true
0 otherwise



- Remember why?

let's use a_k to represent value of A on the k th example

MLE estimate of $\theta_{s|ij}$ from fully observed data

- Maximum likelihood estimate

$$\theta \leftarrow \arg \max_{\theta} \log P(\text{data}|\theta)$$

- Our case:


$$P(\text{data}|\theta) = \prod_{k=1}^K P(f_k, a_k, s_k, h_k, n_k)$$

$$P(\text{data}|\theta) = \prod_{k=1}^K P(f_k)P(a_k)P(s_k|f_k a_k)P(h_k|s_k)P(n_k|s_k)$$

$$\log P(\text{data}|\theta) = \sum_{k=1}^K \log P(f_k) + \log P(a_k) + \log P(s_k|f_k a_k) + \log P(h_k|s_k) + \log P(n_k|s_k)$$

$$\frac{\partial \log P(\text{data}|\theta)}{\partial \theta_{s|ij}} = \sum_{k=1}^K \frac{\partial \log P(s_k|f_k a_k)}{\partial \theta_{s|ij}}$$

$$\theta_{s|ij} = \frac{\sum_{k=1}^K \delta(f_k = i, a_k = j, s_k = 1)}{\sum_{k=1}^K \delta(f_k = i, a_k = j)}$$




MLE for $\theta_{s|ij} = P(S = 1|F = i, A = j)$ from fully observed data

- Maximum likelihood estimate

$$\theta \leftarrow \arg \max_{\theta} \log P(\text{data}|\theta)$$

$$\theta_{s|ij} = \frac{\sum_{k=1}^K \delta(f_k = i, a_k = j, s_k = 1)}{\sum_{k=1}^K \delta(f_k = i, a_k = j)}$$




like flipping coin $\sum_{k=1}^K \delta(f_k = i, a_k = j)$ times to see
how often $s_k = 1$

MAP for $\theta_{s|ij} = P(S = 1|F = i, A = j)$ from fully observed data

- Maximum likelihood estimate

$$\theta \leftarrow \arg \max_{\theta} \log P(\text{data}|\theta)$$

$$\theta_{s|ij} = \frac{\sum_{k=1}^K \delta(f_k=i, a_k=j, s_k=1)}{\sum_{k=1}^K \delta(f_k=i, a_k=j)}$$



- MAP estimate

$$\theta \leftarrow \arg \max_{\theta} \log P(\theta|\text{data}) = \arg \max_{\theta} \log [P(\text{data}|\theta)P(\theta)]$$

If assume prior $P(\theta_{s|ij}) = \text{Beta}(\beta_1, \beta_0) = \frac{1}{B(\beta_1, \beta_0)} \theta_{s|ij}^{\beta_1-1} (1 - \theta_{s|ij})^{\beta_0-1}$

$$\theta_{s|ij} = \frac{(\beta_1-1) + \sum_{k=1}^K \delta(f_k=i, a_k=j, s_k=1)}{(\beta_1-1) + (\beta_0-1) + \sum_{k=1}^K \delta(f_k=i, a_k=j)}$$




like coin flipping, including hallucinated examples

Estimate θ from partly observed data

- What if FAHN observed, but not S?
- Can't calculate MLE

$$\theta \leftarrow \arg \max_{\theta} \log \prod_k P(f_k, a_k, s_k, h_k, n_k | \theta)$$




- Let X be all *observed* variable values (over all examples)
- Let Z be all *unobserved* variable values
- Can't calculate MLE:
$$\theta \leftarrow \arg \max_{\theta} \log P(X, Z | \theta)$$

- WHAT TO DO?

Estimate θ from partly observed data

- What if FAHN observed, but not S?
- Can't calculate MLE

$$\theta \leftarrow \arg \max_{\theta} \log \prod_k P(f_k, a_k, s_k, h_k, n_k | \theta)$$



- Let X be all *observed* variable values (over all examples)
- Let Z be all *unobserved* variable values
- Can't calculate MLE:

$$\theta \leftarrow \arg \max_{\theta} \log P(X, Z | \theta)$$

- EM seeks* the estimate:

$$\theta \leftarrow \arg \max_{\theta} E_{Z|X, \theta} [\log P(X, Z | \theta)]$$


* EM guaranteed to find local maximum

Expected value

$$E_{P(X)}[f(X)] = \sum_x P(X = x) f(x)$$

- EM seeks estimate:

$$\theta \leftarrow \arg \max_{\theta} E_{Z|X,\theta}[\log P(X, Z|\theta)]$$




- here, observed $X=\{F,A,H,N\}$, unobserved $Z=\{S\}$

$$\log P(X, Z|\theta) = \sum_{k=1}^K \log P(f_k) + \log P(a_k) + \log P(s_k|f_k a_k) + \log P(h_k|s_k) + \log P(n_k|s_k)$$

$$E_{P(Z|X,\theta)} \log P(X, Z|\theta) = \sum_{k=1}^K \sum_{i=0}^1 P(s_k = i | f_k, a_k, h_k, n_k) [\log P(f_k) + \log P(a_k) + \log P(s_k | f_k a_k) + \log P(h_k | s_k) + \log P(n_k | s_k)]$$

let's use a_k to represent value of A on the k th example

EM Algorithm - Informally



EM is a general procedure for learning from partly observed data

Given observed variables X , unobserved Z ($X=\{F,A,H,N\}$, $Z=\{S\}$)

Begin with arbitrary choice for parameters θ

Iterate until convergence:

- E Step: use X , θ to estimate the unobserved Z values
- M Step: use X values and estimated Z values to derive a better θ

Guaranteed to find local maximum.

Each iteration increases $E_{P(Z|X,\theta)}[\log P(X, Z|\theta')]$

EM Algorithm - Precisely

EM is a general procedure for learning from partly observed data

Given observed variables X , unobserved Z ($X=\{F,A,H,N\}$, $Z=\{S\}$)

Define $Q(\theta'|\theta) = E_{P(Z|X,\theta)}[\log P(X, Z|\theta')]$

current *M step new*

Iterate until convergence:

- E Step: Use X and current θ to calculate $P(Z|X,\theta)$
- M Step: Replace current θ by


$$\theta \leftarrow \arg \max_{\theta'} Q(\theta'|\theta)$$

Guaranteed **to** find local maximum.

Each iteration increases $E_{P(Z|X,\theta)}[\log P(X, Z|\theta')]$

E Step: Use X, θ , to Calculate $P(Z|X, \theta)$

observed $X=\{F,A,H,N\}$,
unobserved $Z=\{S\}$




How? Bayes net inference problem.

$$P(S_k = 1 | f_k a_k h_k n_k, \theta) =$$

let's use a_k to represent value of A on the k th example

E Step: Use X, θ , to Calculate $P(Z|X, \theta)$

observed $X=\{F,A,H,N\}$,
unobserved $Z=\{S\}$



How? Bayes net inference problem.


$$P(S_k = 1 | f_k a_k h_k n_k, \theta) =$$

$$P(S_k = 1 | f_k a_k h_k n_k, \theta) = \frac{P(S_k = 1, f_k a_k h_k n_k | \theta)}{P(S_k = 1, f_k a_k h_k n_k | \theta) + P(S_k = 0, f_k a_k h_k n_k | \theta)}$$

let's use a_k to represent value of A on the k th example

EM and estimating $\theta_{s|ij}$

observed $X = \{F, A, H, N\}$, unobserved $Z = \{S\}$



E step: Calculate $P(Z_k|X_k; \theta)$ for each training example, k


$$P(S_k = 1 | f_k a_k h_k n_k, \theta) = E[s_k] = \frac{P(S_k = 1, f_k a_k h_k n_k | \theta)}{P(S_k = 1, f_k a_k h_k n_k | \theta) + P(S_k = 0, f_k a_k h_k n_k | \theta)}$$

M step: update all relevant parameters. For example:

$$\theta_{s|ij} \leftarrow \frac{\sum_{k=1}^K \delta(f_k = i, a_k = j) E[s_k]}{\sum_{k=1}^K \delta(f_k = i, a_k = j)}$$

Recall MLE was: $\theta_{s|ij} = \frac{\sum_{k=1}^K \delta(f_k = i, a_k = j, s_k = 1)}{\sum_{k=1}^K \delta(f_k = i, a_k = j)}$

EM and estimating θ



More generally,

Given observed set X , unobserved set Z of boolean values

E step: Calculate for each training example, k

the expected value of each unobserved variable in
each training example

M step:


Calculate θ similar to MLE estimates, but
replacing each count by its expected count

$$\delta(Z = 1) \rightarrow E_{Z|X,\theta}[Z]$$

$$\delta(Z = 0) \rightarrow (1 - E_{Z|X,\theta}[Z])$$


Using Unlabeled Data to Help Train Naïve Bayes Classifier

Learn $P(Y|X)$




Y	X1	X2	X3	X4
1	0	0	1	1
0	0	1	0	0
0	0	0	1	0
?	0	1	1	0
?	0	1	0	1

E step: Calculate for each training example, k
the expected value of each unobserved variable



Y	X1	X2	X3	X4
1	0	0	1	1
0	0	1	0	0
0	0	0	1	0
?	0	1	1	0
?	0	1	0	1

EM and estimating θ



Given observed set X , unobserved set Y of boolean values

E step: Calculate for each training example, k


the expected value of each unobserved variable Y

$$E_{P(Y|X_1 \dots X_N)}[y(k)] = P(y(k) = 1 | x_1(k), \dots, x_N(k); \theta) = \frac{P(y(k) = 1) \prod_i P(x_i(k) | y(k) = 1)}{\sum_{j=0}^1 P(y(k) = j) \prod_i P(x_i(k) | y(k) = j)}$$

M step: Calculate estimates similar to MLE, but replacing each count by its expected count

let's use $y(k)$ to indicate value of Y on k th example

EM and estimating θ



Given observed set X , unobserved set Y of boolean values

E step: Calculate for each training example, k

the expected value of each unobserved variable Y

$$E_{P(Y|X_1 \dots X_N)}[y(k)] = P(y(k) = 1 | x_1(k), \dots, x_N(k); \theta) = \frac{P(y(k) = 1) \prod_i P(x_i(k) | y(k) = 1)}{\sum_{j=0}^1 P(y(k) = j) \prod_i P(x_i(k) | y(k) = j)}$$

M step: Calculate estimates similar to MLE, but replacing each count by its expected count

$$\theta_{ij|m} = \hat{P}(X_i = j | Y = m) = \frac{\sum_k P(y(k) = m | x_1(k) \dots x_N(k)) \delta(x_i(k) = j)}{\sum_k P(y(k) = m | x_1(k) \dots x_N(k))}$$

MLE would be:
$$\hat{P}(X_i = j | Y = m) = \frac{\sum_k \delta((y(k) = m) \wedge (x_i(k) = j))}{\sum_k \delta(y(k) = m)}$$

-
- **Inputs:** Collections \mathcal{D}^l of labeled documents and \mathcal{D}^u of unlabeled documents.
 - Build an initial naive Bayes classifier, $\hat{\theta}$, from the labeled documents, \mathcal{D}^l , only. Use maximum a posteriori parameter estimation to find $\hat{\theta} = \arg \max_{\theta} P(\mathcal{D}|\theta)P(\theta)$ (see Equations 5 and 6).
 - Loop while classifier parameters improve, as measured by the change in $l_c(\theta|\mathcal{D}; \mathbf{z})$ (the complete log probability of the labeled and unlabeled data)
 - **(E-step)** Use the current classifier, $\hat{\theta}$, to estimate component membership of each unlabeled document, *i.e.*, the probability that each mixture component (and class) generated each document, $P(c_j|d_i; \hat{\theta})$ (see Equation 7).
 - **(M-step)** Re-estimate the classifier, $\hat{\theta}$, given the estimated component membership of each document. Use maximum a posteriori parameter estimation to find $\hat{\theta} = \arg \max_{\theta} P(\mathcal{D}|\theta)P(\theta)$ (see Equations 5 and 6).
 - **Output:** A classifier, $\hat{\theta}$, that takes an unlabeled document and predicts a class label.

From [Nigam et al., 2000]



Experimental Evaluation

- Newsgroup postings
 - 20 newsgroups, 1000/group
- Web page classification
 - student, faculty, course, project
 - 4199 web pages
- Reuters newswire articles
 - 12,902 articles
 - 90 topics categories

20 Newsgroups




Table 3. Lists of the words most predictive of the **course** class in the WebKB data set, as they change over iterations of EM for a specific trial. By the second iteration of EM, many common **course**-related words appear. The symbol *D* indicates an arbitrary digit.

Iteration 0		Iteration 1	Iteration 2
intelligence	word w ranked by $P(w Y=\text{course}) / P(w Y \neq \text{course})$	<i>DD</i>	<i>D</i>
<i>DD</i>		<i>D</i>	<i>DD</i>
artificial		lecture	lecture
understanding		cc	cc
<i>DDw</i>		<i>D</i> *	<i>DD:DD</i>
dist		<i>DD:DD</i>	due
identical		handout	<i>D</i> *
rus		due	homework
arrange		problem	assignment
games		set	handout
dartmouth	Using one labeled example per class	tay	set
natural		<i>DDam</i>	hw
cognitive		yurttas	exam
logic		homework	problem
proving		kfoury	<i>DDam</i>
prolog		sec	postscript
knowledge		postscript	solution
human		exam	quiz
representation		solution	chapter
field		assaf	ascii

What you should know about EM

- For learning from partly unobserved data
- MLE of $\theta = \arg \max_{\theta} \log P(\text{data}|\theta)$
- EM estimate: $\theta = \arg \max_{\theta} E_{Z|X,\theta}[\log P(X, Z|\theta)]$
Where X is observed part of data, Z is unobserved
- EM for training Bayes networks
- Can also develop MAP version of EM
- Can also derive your own EM algorithm for your own problem
 - write out expression for $E_{Z|X,\theta}[\log P(X, Z|\theta)]$
 - E step: for each training example X^k , calculate $P(Z^k | X^k, \theta)$
 - M step: chose new θ to maximize $E_{Z|X,\theta}[\log P(X, Z|\theta)]$

Learning Bayes Net Structure

How can we learn Bayes Net graph structure?

In general case, open problem

- can require lots of data (else high risk of overfitting)
- can use Bayesian priors, or other kinds of prior assumptions about graph structure to constrain search

One key result:

- Chow-Liu algorithm: finds “best” tree-structured network

- What’s the cost?

- support network $KL(P(\mathbf{X}) \parallel T(\mathbf{X})) \equiv \sum_k P(\mathbf{X} = k) \log \frac{P(\mathbf{X} = k)}{T(\mathbf{X} = k)}$

- Chow-Liu minimizes Kullback-Leibler divergence:

Chow-Liu Algorithm

Key result: To minimize $KL(P \parallel T)$ over possible tree networks T representing true P , it suffices to find the tree network T that maximizes the sum of mutual informations over its edges

Mutual information for an edge between variable A and B :

$$I(A, B) = \sum_a \sum_b P(a, b) \log \frac{P(a, b)}{P(a)P(b)}$$

This works because for tree networks with nodes $\mathbf{X} \equiv \langle X_1 \dots X_n \rangle$

$$\begin{aligned} KL(P(\mathbf{X}) \parallel T(\mathbf{X})) &\equiv \sum_k P(\mathbf{X} = k) \log \frac{P(\mathbf{X} = k)}{T(\mathbf{X} = k)} \\ &= - \sum_i I(X_i, Pa(X_i)) + \sum_i H(X_i) - H(X_1 \dots X_n) \end{aligned}$$

Chow-Liu Algorithm

1. for each pair of variables A,B, use data to estimate $P(A,B)$, $P(A)$, and $P(B)$






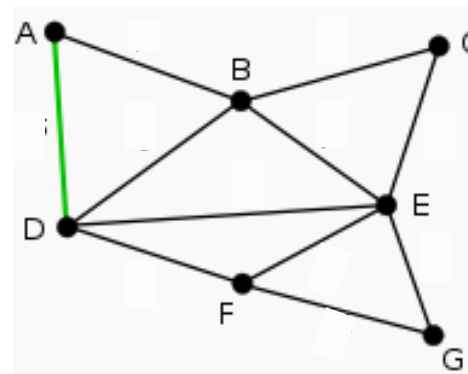
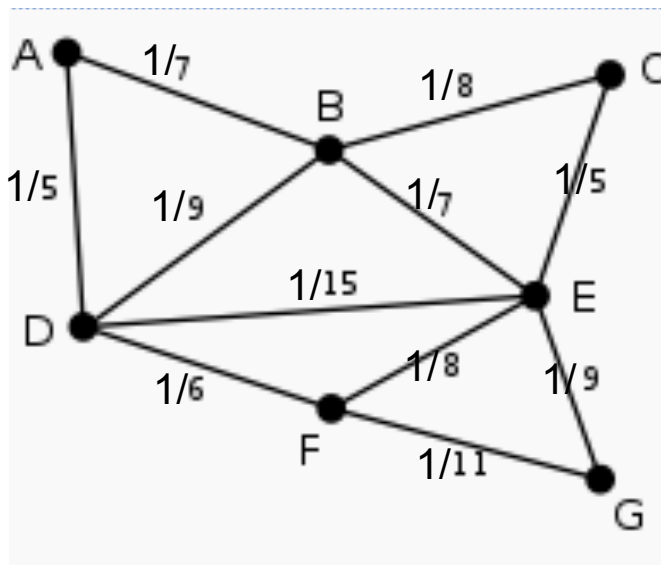
2. for each pair A, B calculate mutual information

$$I(A, B) = \sum_a \sum_b P(a, b) \log \frac{P(a, b)}{P(a)P(b)} \quad \text{🗨️}$$

3. calculate the maximum spanning tree over the set of variables, using edge weights $I(A,B)$
(given N vars, this costs only $O(N^2)$ time)
4. add arrows to edges to form a directed-acyclic graph
5. learn the CPD's for this graph

Chow-Liu algorithm example


Greedy Algorithm to find Max-Spanning Tree



[courtesy A. Singh, C. Guestrin]

Tree Augmented Naïve Bayes

[Nir Friedman et al., 1997]



Bayes Nets – What You Should Know

- Representation
 - Bayes nets represent joint distribution as a DAG + Conditional Distributions
 - D-separation lets us decode conditional independence assumptions
- Inference
 - NP-hard in general
 - For some graphs, closed form inference is feasible
 - Approximate methods too, e.g., Monte Carlo methods, ...
- Learning
 - Easy for known graph, fully observed data (MLE's, MAP est.)
 - EM for partly observed data, known graph
 - Learning graph structure: Chow-Liu for tree-structured networks
 - Hardest when graph unknown, data incompletely observed