

# Machine Learning 10-601

Tom M. Mitchell  
Machine Learning Department  
Carnegie Mellon University

August 30, 2017

## Today:

- Decision trees
- Overfitting
- The Big Picture

## Coming soon

- Probabilistic learning
- MLE, MAP estimates

## Readings:

- Decision trees, overfitting
- Mitchell, Chapter 3

## Probabilistic learning

- [Estimating Probabilities \[Mitchell\]](#)
- [Andrew Moore's online probability tutorial](#)

## Function Approximation:

### Problem Setting:

- Set of possible instances  $X$
- Unknown target function  $f: X \rightarrow Y$
- Set of function hypotheses  $H = \{ h \mid h: X \rightarrow Y \}$

### Input:

- Training examples  $\{ \langle x^{(i)}, y^{(i)} \rangle \}$  of unknown target function  $f$

### Output:

- Hypothesis  $h \in H$  that best approximates target function  $f$

## Function Approximation: Decision Tree Learning

### Problem Setting:

- Set of possible instances  $X$ 
  - each instance  $x$  in  $X$  is a feature vector  
 $x = \langle x_1, x_2 \dots x_n \rangle$
- Unknown target function  $f: X \rightarrow Y$ 
  - $Y$  is discrete valued
- Set of function hypotheses  $H = \{ h \mid h: X \rightarrow Y \}$ 
  - each hypothesis  $h$  is a decision tree


### Input:

- Training examples  $\{ \langle x^{(i)}, y^{(i)} \rangle \}$  of unknown target function  $f$

### Output:

- Hypothesis  $h \in H$  that best approximates target function  $f$


## Function approximation as Search for the best hypothesis



- ID3 performs heuristic search through space of decision trees

## Function Approximation: The Big Picture

## Which Tree Should We Output?



- ID3 performs heuristic search through space of decision trees
- It stops at smallest acceptable tree. Why?

Occam's razor: prefer the simplest hypothesis that fits the data

## Why Prefer Short Hypotheses? (Occam's Razor)

Arguments in favor:



Arguments opposed:



## Why Prefer Short Hypotheses? (Occam's Razor)

Argument in favor:

- Fewer short hypotheses than long ones
- a short hypothesis that fits the data is less likely to be a statistical coincidence

Argument opposed:

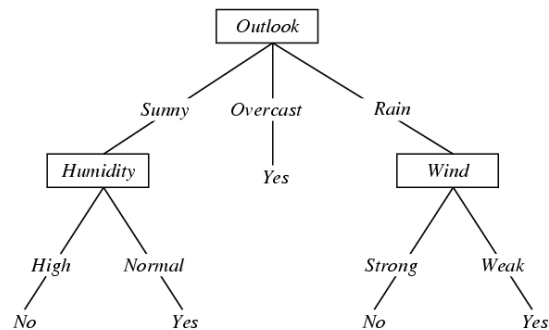
- Also fewer hypotheses containing a prime number of nodes and attributes beginning with "Z"
- What's so special about "short" hypotheses, instead of "prime number of nodes and edges"?

## Overfitting in Decision Trees

Consider adding noisy training example #15:

*Sunny, Mild, Normal, Strong, PlayTennis=No*

What effect on earlier tree?



## Overfitting

Consider a hypothesis  $h$  and its

- Error rate over training data:  $error_{train}(h)$
- True error rate over all data:  $error_{true}(h)$

## Overfitting

Consider a hypothesis  $h$  and its

- Error rate over training data:  $error_{train}(h)$
- True error rate over all data:  $error_{true}(h)$

We say  $h$  overfits the training data if


$$error_{true}(h) > error_{train}(h)$$

Amount of overfitting =

$$error_{true}(h) - error_{train}(h)$$

## Overfitting in Decision Tree Learning

---



## Avoiding Overfitting

---

How can we avoid overfitting?

- stop growing when data split not statistically significant
- grow full tree, then post-prune

## How Can We Avoid Overfitting?

1. stop growing tree when data split is not statistically significant
2. grow full tree, then post-prune
3. learn a collection of trees (decision forest) by randomizing training, then have them vote

## Reduced-Error Pruning

---

Split data into *training* and *validation* set


Learn a tree that classifies *training* set correctly

Do until further pruning is harmful:

1. For each non-leaf node, evaluate impact on *validation* set of converting it to a leaf node
  2. Greedily select the node that would most improve *validation* set accuracy, and convert it to a leaf
- this produces smallest version of most accurate (over the *validation* set) subtree



## Effect of Reduced-Error Pruning



## Decision Forests

Key idea:

1. learn a collection of many trees
2. classify by taking a weighted vote of the trees

Empirically successful. Widely used in industry.

- human pose recognition in Microsoft kinect
- medical imaging – cortical parcellation
- classify disease from gene expression data

How to train different trees

1. Train on different random subsets of data
2. Randomize the choice of decision nodes

## Decision Forests

Key idea:

1. learn a collection of many trees
2. classify by taking a weighted vote of the trees

more to come

Em

- h
  - m
  - c
- later lecture on boosting and ensemble methods...



How

1. Train on different random subsets of data
2. Randomize the choice of decision nodes

## You should know:

- Well posed function approximation problems:
  - Instance space,  $X$
  - Sample of labeled training data  $\{ \langle x^{(i)}, y^{(i)} \rangle \}$
  - Hypothesis space,  $H = \{ f: X \rightarrow Y \}$
- Learning is a search/optimization problem over  $H$ 
  - Various objective functions to define the goal
    - minimize training error (0-1 loss)
    - minimize validation error (0-1 loss)
    - among hypotheses that minimize error, select smallest (?)
- Decision tree learning
  - Greedy top-down learning of decision trees (ID3, C4.5, ...)
  - Overfitting and post-pruning
  - Extensions... to continuous values, probabilistic classification
  - Widely used commercially: decision *forests*

## Further Reading...



## Extra slides

extensions to decision tree learning

## Rule Post-Pruning


---

1. Convert tree to equivalent set of rules
2. Prune each rule independently of others
3. Sort final rules into desired sequence for use

frequently used method (e.g., C4.5)

## Converting A Tree to Rules

---



## Unknown Attribute Values

---

What if some examples missing values of  $A$ ?

Use training example anyway, sort through tree

- If node  $n$  tests  $A$ , assign most common value of  $A$  among other examples sorted to node  $n$
- assign most common value of  $A$  among other examples with same target value
- assign probability  $p_i$  to each possible value  $v_i$  of  $A$ 
  - assign fraction  $p_i$  of example to each descendant in tree

Classify new examples in same fashion

## Questions to think about (1)

- Consider target function  $f: \langle x_1, x_2 \rangle \rightarrow y$ , where  $x_1$  and  $x_2$  are real-valued,  $y$  is boolean. What is the set of decision surfaces describable with decision trees that use each attribute at most once?

## Questions to think about (2)

- ID3 and C4.5 are heuristic algorithms that search through the space of decision trees. Why not just do an exhaustive search?

## Questions to think about (3)

- Why use Information Gain to select attributes in decision trees? What other criteria seem reasonable, and what are the tradeoffs in making this choice?

probabilistic function approximation:

instead of  $F: X \rightarrow Y$ ,  
learn  $P(Y | X)$

## Random Variables

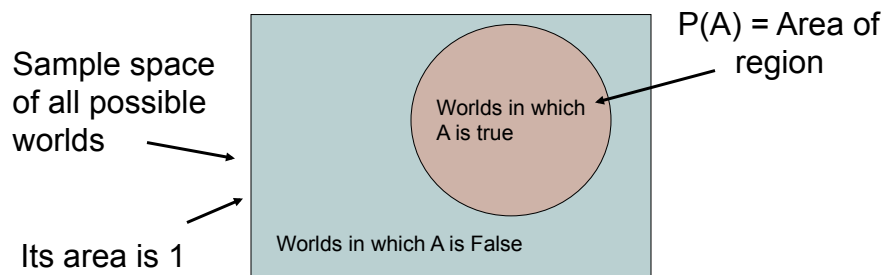
- Informally,  $A$  is a random variable if
  - $A$  denotes something about which we are uncertain
  - perhaps the outcome of a randomized experiment
- Examples
  - $A$  = True if a randomly drawn person from our class is female
  - $A$  = The hometown of a randomly drawn person from our class
  - $A$  = True if two randomly drawn persons from our class have same birthday
- Define  $P(A)$  as “the fraction of possible worlds in which  $A$  is true” or “the fraction of times  $A$  holds, in repeated runs of the random experiment”
  - the set of possible worlds is called the sample space,  $S$
  - A random variable  $A$  is a function defined over  $S$   
 $A: S \rightarrow \{0,1\}$

## A little formalism

More formally, we have

- a sample space  $S$  (e.g., set of students in our class)
  - aka the set of possible worlds
- a random variable is a function defined over the sample space
  - Gender:  $S \rightarrow \{m, f\}$
  - Height:  $S \rightarrow \text{Reals}$
- an event is a subset of  $S$ 
  - e.g., the subset of  $S$  for which Gender=f
  - e.g., the subset of  $S$  for which (Gender=m) AND (Height > 2m)
- we're often interested in probabilities of specific events
- and of specific events conditioned on other specific events


## Visualizing A





## Elementary Probability in Pictures

- $P(\sim A) + P(A) = 1$



## The Axioms of Probability

- $0 \leq P(A) \leq 1$
- $P(\text{True}) = 1$
- $P(\text{False}) = 0$
- $P(A \text{ or } B) = P(A) + P(B) - P(A \text{ and } B)$

[di Finetti 1931]:

when gambling based on “uncertainty formalism A” you can be exploited by an opponent

iff

your uncertainty formalism A violates these axioms

## A useful theorem


- Axioms:  $0 \leq P(A) \leq 1$ ,  $P(\text{True}) = 1$ ,  $P(\text{False}) = 0$ ,  
 $P(A \vee B) = P(A) + P(B) - P(A \wedge B)$

$$\rightarrow P(A) = P(A \wedge B) + P(A \wedge \sim B)$$

*prove this yourself*


## Elementary Probability in Pictures

- $P(A) = P(A \wedge B) + P(A \wedge \sim B)$



## Definition of Conditional Probability

$$P(A|B) = \frac{P(A \cap B)}{P(B)}$$



## Definition of Conditional Probability


$$P(A|B) = \frac{P(A \cap B)}{P(B)}$$

## Corollary: The Chain Rule

$$P(A \cap B) = P(A|B) P(B)$$


## Bayes Rule

- let's write 2 expressions for  $P(A \wedge B)$



## Bayes Rule

- let's write 2 expressions for  $P(A \wedge B)$



$$P(A \wedge B) = P(A|B)P(B) = P(B|A) P(B)$$

implies: 
$$P(A|B) = \frac{P(B|A) * P(A)}{P(B)}$$

$$P(A|B) = \frac{P(B|A) * P(A)}{P(B)} \quad \text{Bayes' rule}$$

we call  $P(A)$  the “prior”

and  $P(A|B)$  the “posterior”



**Bayes, Thomas (1763)** An essay towards solving a problem in the doctrine of chances. *Philosophical Transactions of the Royal Society of London*, **53:370-418**

## Other Forms of Bayes Rule

$$P(A|B) = \frac{P(B|A)P(A)}{P(B|A)P(A) + P(B|\sim A)P(\sim A)}$$

$$P(A|B \wedge X) = \frac{P(B|A \wedge X)P(A \wedge X)}{P(B \wedge X)}$$

## Applying Bayes Rule

$$P(A|B) = \frac{P(B|A)P(A)}{P(B|A)P(A) + P(B|\sim A)P(\sim A)}$$

A = you have the flu, B = you just coughed

Assume:

$$P(A) = 0.05$$

$$P(B|A) = 0.80$$

$$P(B|\sim A) = 0.2$$

what is  $P(\text{flu} | \text{cough}) = P(A|B)$ ?

## The Awesome Joint Probability Distribution $P(X_1, X_2, \dots, X_N)$

from which we can calculate

$$P(X_1|X_2 \dots X_N),$$

and every other probability we desire  
over subsets of  $X_1 \dots X_N$

## The Joint Distribution

*Example: Boolean variables A, B, C*

Recipe for making a joint distribution of M variables:

## The Joint Distribution

*Example: Boolean variables A, B, C*

Recipe for making a joint distribution of M variables:

1. Make a table listing all combinations of values of your variables (if there are M Boolean variables then the table will have  $2^M$  rows).

A	B	C
0	0	0
0	0	1
0	1	0
0	1	1
1	0	0
1	0	1
1	1	0
1	1	1

## The Joint Distribution

*Example: Boolean variables A, B, C*

Recipe for making a joint distribution of M variables:

1. Make a table listing all combinations of values of your variables (if there are M Boolean variables then the table will have  $2^M$  rows).
2. For each combination of values, say how probable it is.


A	B	C	Prob
0	0	0	0.30
0	0	1	0.05
0	1	0	0.10
0	1	1	0.05
1	0	0	0.05
1	0	1	0.10
1	1	0	0.25
1	1	1	0.10

## The Joint Distribution

Recipe for making a joint distribution of M variables:









1. Make a table listing all combinations of values of your variables (if there are M Boolean variables then the table will have  $2^M$  rows).
2. For each combination of values, say how probable it is.
3. If you subscribe to the axioms of probability, those numbers must sum to 1.

A	B	C	Prob
0	0	0	0.30
0	0	1	0.05
0	1	0	0.10
0	1	1	0.05
1	0	0	0.05
1	0	1	0.10
1	1	0	0.25
1	1	1	0.10













## Using the Joint Distribution

gender	hours_worked	wealth	
Female	v0:40.5-	poor	0.253122 
		rich	0.0245895 
	v1:40.5+	poor	0.0421768 
		rich	0.0116293 
Male	v0:40.5-	poor	0.331313 
		rich	0.0971295 
	v1:40.5+	poor	0.134106 
		rich	0.105933 

One you have the JD  
you can ask for the  
probability of **any** logical  
expression involving  
these variables

$$P(E) = \sum_{\text{rows matching } E} P(\text{row})$$

## Using the Joint

gender	hours_worked	wealth	
Female	v0:40.5-	poor	0.253122 
		rich	0.0245895 
	v1:40.5+	poor	0.0421768 
		rich	0.0116293 
Male	v0:40.5-	poor	0.331313 
		rich	0.0971295 
	v1:40.5+	poor	0.134106 
		rich	0.105933 

$$P(\text{Poor Male}) = 0.4654$$

$$P(E) = \sum_{\text{rows matching } E} P(\text{row})$$

## Using the Joint

gender	hours_worked	wealth	
Female	v0:40.5-	poor	0.253122
		rich	0.0245895
	v1:40.5+	poor	0.0421768
		rich	0.0116293
Male	v0:40.5-	poor	0.331313
		rich	0.0971295
	v1:40.5+	poor	0.134106
		rich	0.105933

$$P(\text{Poor}) = 0.7604$$

$$P(E) = \sum_{\text{rows matching } E} P(\text{row})$$









## Inference with the Joint

gender	hours_worked	wealth	
Female	v0:40.5-	poor	0.253122
		rich	0.0245895
	v1:40.5+	poor	0.0421768
		rich	0.0116293
Male	v0:40.5-	poor	0.331313
		rich	0.0971295
	v1:40.5+	poor	0.134106
		rich	0.105933

$$P(E_1 | E_2) = \frac{P(E_1 \wedge E_2)}{P(E_2)} = \frac{\sum_{\text{rows matching } E_1 \text{ and } E_2} P(\text{row})}{\sum_{\text{rows matching } E_2} P(\text{row})}$$

$$P(\text{Male} | \text{Poor}) = 0.4654 / 0.7604 = 0.612$$

## Learning and the Joint Distribution

gender	hours_worked	wealth	
Female	v0:40.5-	poor	0.253122 
		rich	0.0245895 
	v1:40.5+	poor	0.0421768 
		rich	0.0116293 
Male	v0:40.5-	poor	0.331313 
		rich	0.0971295 
	v1:40.5+	poor	0.134106 
		rich	0.105933 

Suppose we want to learn the function  $f: \langle G, H \rangle \rightarrow W$

Equivalently,  $P(W | G, H)$

Solution: learn joint distribution from data, calculate  $P(W | G, H)$

e.g.,  $P(W=\text{rich} | G = \text{female}, H = 40.5- ) =$

sounds like the solution to  
learning  $F: X \rightarrow Y$ ,  
or  $P(Y | X)$ .

Are we done?

sounds like the solution to  
learning  $F: X \rightarrow Y$ ,  
or  $P(Y | X)$ .

Main problem: learning  $P(Y|X)$   
can require more data than we have

consider learning Joint Dist. with 100 attributes

# of rows in this table?

# of people on earth?

## What to do?

1. Be smart about how we estimate probabilities from sparse data
  - maximum likelihood estimates
  - maximum a posteriori estimates
2. Be smart about how to represent joint distributions
  - Bayes networks, graphical models

## 1. Be smart about how we estimate probabilities

### Estimating Probability of Heads



- I show you the above coin  $X$ , and ask you to estimate the probability that it will turn up heads ( $X=1$ ) or tails ( $X=0$ )
- You flip it repeatedly, observing
  - it turns up heads  $\alpha_1$  times
  - it turns up tails  $\alpha_0$  times
- Your estimate for  $\hat{\theta} = \hat{P}(X = 1)$  is ...?

## Estimating Probability of Heads



- I show you the above coin  $X$ , and ask you to estimate the probability that it will turn up heads ( $X=1$ ) or tails ( $X=0$ )
- You flip it repeatedly, observing
  - it turns up heads  $\alpha_1$  times
  - it turns up tails  $\alpha_0$  times

Algorithm 1 (MLE):  $\hat{\theta} = \hat{P}(X = 1) = \frac{\alpha_1}{\alpha_1 + \alpha_0}$

## Estimating $\theta = P(X=1)$



Test A:

100 flips: 51 Heads, 49 Tails

Test B:

3 flips: 2 Heads, 1 Tails

## Estimating Probability of Heads



When data sparse, might bring in prior assumptions to bias our estimate

- e.g., represent priors by “hallucinating”  $\gamma_1$  heads, and  $\gamma_0$  tails, to complement sparse observations

$$\text{Alg 2 (MAP): } \hat{\theta} = \hat{P}(X = 1) = \frac{(\alpha_1 + \gamma_1)}{(\alpha_1 + \gamma_1) + (\alpha_0 + \gamma_0)}$$

## Estimating Probability of Heads





When data sparse, might bring in prior assumptions to bias our estimate

- e.g., represent priors by “hallucinating”  $\gamma_1$  heads, and  $\gamma_0$  tails, to complement sparse observations

$$\text{Alg 2 (MAP): } \hat{\theta} = \hat{P}(X = 1) = \frac{(\alpha_1 + \gamma_1)}{(\alpha_1 + \gamma_1) + (\alpha_0 + \gamma_0)}$$

Consider  $\gamma_1 = 1$   $\gamma_0 = 1$   
versus  $\gamma_1 = 1000$   $\gamma_0 = 1000$   
versus  $\gamma_1 = 500$   $\gamma_0 = 1500$





## Principles for Estimating Probabilities

- Maximum Likelihood Estimate (MLE): choose  $\theta$  that maximizes probability of observed data  $\mathcal{D}$

$$\hat{\theta} = \arg \max_{\theta} P(\mathcal{D} | \theta)$$

- Maximum a Posteriori (MAP) estimate: choose  $\theta$  that is most probable given prior probability and observed data

$$\begin{aligned}\hat{\theta} &= \arg \max_{\theta} P(\theta | \mathcal{D}) \\ &= \arg \max_{\theta} \frac{P(\mathcal{D} | \theta)P(\theta)}{P(\mathcal{D})} \\ &= \arg \max_{\theta} P(\mathcal{D} | \theta)P(\theta)\end{aligned}$$

## Principles for Estimating Probabilities

Principle 1 (maximum likelihood):

- choose parameters  $\theta$  that maximize **P(data |  $\theta$ )**
- result in our case:  $\hat{\theta}^{MLE} = \frac{\alpha_1}{\alpha_1 + \alpha_0}$

Principle 2 (maximum a posteriori probability):

- choose parameters  $\theta$  that maximize **P( $\theta$  | data)**
- result in our case:

$$\hat{\theta}^{MAP} = \frac{\alpha_1 + \text{\#hallucinated\_1s}}{(\alpha_1 + \text{\#hallucinated\_1s}) + (\alpha_0 + \text{\#hallucinated\_0s})}$$

## Maximum Likelihood Estimation

given data D, choose  $\theta$  that maximizes  $P(D | \theta)$

Data D:

$$P(D|\theta) =$$



X=1

X=0

$$P(X=1) = \theta$$

$$P(X=0) = 1-\theta$$

(Bernoulli)

## Maximum Likelihood Estimation

given data D, choose  $\theta$  that maximizes  $P(D | \theta)$

Data D: < 1 0 0 1 1 >

$$\begin{aligned} P(D|\theta) &= \theta \cdot (1 - \theta) \cdot (1 - \theta) \cdot \theta \cdot \theta \\ &= \theta^{\alpha_1} \cdot (1 - \theta)^{\alpha_0} \end{aligned}$$



X=1

X=0

$$P(X=1) = \theta$$

$$P(X=0) = 1-\theta$$

(Bernoulli)

Flips are independent, identically distributed 1's and 0's,  
producing  $\alpha_1$  1's, and  $\alpha_0$  0's

$$\begin{aligned} \text{Now solve for: } \hat{\theta}^{MLE} &= \arg \max_{\theta} P(D|\theta) \\ &= \arg \max_{\theta} P(\alpha_1, \alpha_0 | \theta) \\ &= \arg \max_{\theta} \theta^{\alpha_1} (1 - \theta)^{\alpha_0} \end{aligned}$$

$$\hat{\theta} = \arg \max_{\theta} \ln P(D|\theta)$$

■ Set derivative to zero:

$$\frac{d}{d\theta} \ln P(D|\theta) = 0$$

$$= \arg \max_{\theta} \ln [\theta^{\alpha_1} (1 - \theta)^{\alpha_0}]$$

hint:  $\frac{\partial \ln \theta}{\partial \theta} = \frac{1}{\theta}$

## Summary: Maximum Likelihood Estimate for Bernoulli random variable



X=1

X=0

$$P(X=1) = \theta$$

$$P(X=0) = 1-\theta$$

(Bernoulli)

- Each flip yields boolean value for  $X$

$$X \sim \text{Bernoulli}: P(X) = \theta^X (1 - \theta)^{(1-X)}$$

- Data set  $D$  of independent, identically distributed (iid) flips produces  $\alpha_1$  ones,  $\alpha_0$  zeros (Binomial)

$$P(D|\theta) = P(\alpha_1, \alpha_0|\theta) = \theta^{\alpha_1} (1 - \theta)^{\alpha_0}$$

$$\hat{\theta}^{MLE} = \operatorname{argmax}_{\theta} P(D|\theta) = \frac{\alpha_1}{\alpha_1 + \alpha_0}$$

## Principles for Estimating Probabilities

Principle 1 (maximum likelihood):

- choose parameters  $\theta$  that maximize  $P(\text{data} \mid \theta)$

Principle 2 (maximum a posteriori prob.):

- choose parameters  $\theta$  that maximize 
$$P(\theta \mid \text{data}) = \frac{P(\text{data} \mid \theta) P(\theta)}{P(\text{data})}$$


## Beta prior distribution : $P(\theta)$

$$P(\theta) = \frac{\theta^{\beta_H-1}(1-\theta)^{\beta_T-1}}{B(\beta_H, \beta_T)} \sim \text{Beta}(\beta_H, \beta_T)$$

- Likelihood function:  $P(\mathcal{D} \mid \theta) = \theta^{\alpha_H}(1-\theta)^{\alpha_T}$
- Posterior:  $P(\theta \mid \mathcal{D}) \propto P(\mathcal{D} \mid \theta)P(\theta)$

## Beta prior distribution – $P(\theta)$

$$P(\theta) = \frac{\theta^{\beta_H-1}(1-\theta)^{\beta_T-1}}{B(\beta_H, \beta_T)} \sim \text{Beta}(\beta_H, \beta_T)$$



### Summary:

#### Maximum a Posteriori (MAP) Estimate for Bernoulli random variable

Likelihood is  $\sim$  Binomial

$$P(\mathcal{D} | \theta) = \theta^{\alpha_H}(1-\theta)^{\alpha_T}$$

If prior is Beta distribution,

$$P(\theta) = \frac{\theta^{\beta_H-1}(1-\theta)^{\beta_T-1}}{B(\beta_H, \beta_T)} \sim \text{Beta}(\beta_H, \beta_T)$$

Then posterior is Beta distribution

$$P(\theta|D) \propto P(D|\theta)P(\theta) \sim \text{Beta}(\alpha_H + \beta_H, \alpha_T + \beta_T)$$

and MAP estimate is therefore

$$\hat{\theta}^{MAP} = \frac{\alpha_H + \beta_H - 1}{(\alpha_H + \beta_H - 1) + (\alpha_T + \beta_T - 1)}$$



$X=1$

$X=0$

$P(X=1) = \theta$

$P(X=0) = 1-\theta$   
(Bernoulli)

### Maximum a Posteriori (MAP) Estimate for random variable with k possible outcomes



Likelihood is  $\sim$  Multinomial( $\theta = \{\theta_1, \theta_2, \dots, \theta_k\}$ )

$$P(\mathcal{D} | \theta) = \theta_1^{\alpha_1} \theta_2^{\alpha_2} \dots \theta_k^{\alpha_k}$$

If prior is Dirichlet distribution,

$$P(\theta) = \frac{\theta_1^{\beta_1-1} \theta_2^{\beta_2-1} \dots \theta_k^{\beta_k-1}}{B(\beta_1, \dots, \beta_k)} \sim \text{Dirichlet}(\beta_1, \dots, \beta_k)$$

Then posterior is Dirichlet distribution

$$P(\theta | \mathcal{D}) \propto P(\mathcal{D} | \theta) P(\theta) \sim \text{Dirichlet}(\alpha_1 + \beta_1, \dots, \alpha_k + \beta_k)$$

and MAP estimate is therefore

$$\hat{\theta}_i^{MAP} = \frac{\alpha_i + \beta_i - 1}{\sum_{j=1}^k (\alpha_j + \beta_j - 1)}$$

### Some terminology

- Likelihood function:  $P(\text{data} | \theta)$
- Prior:  $P(\theta)$
- Posterior:  $P(\theta | \text{data})$
- Conjugate prior:  $P(\theta)$  is the conjugate prior for likelihood function  $P(\text{data} | \theta)$  if the forms of  $P(\theta)$  and  $P(\theta | \text{data})$  are the same.
  - Beta is conjugate prior for Bernoulli, Binomial
  - Dirichlet is conjugate prior for Multinomial

## You should know

- Probability basics
  - random variables, conditional probs, ...
  - Bayes rule
  - Joint probability distributions
  - calculating probabilities from the joint distribution
- Estimating parameters from data
  - maximum likelihood estimates
  - maximum a posteriori estimates
  - distributions – Bernoulli, Binomial, Beta, Dirichlet, ...
  - conjugate priors

Extra slides

## Independent Events

- Definition: two events A and B are *independent* if  $P(A \wedge B) = P(A) * P(B)$
- Intuition: knowing A tells us nothing about the value of B (and vice versa)

Picture “A independent of B”



## Expected values

Given a discrete random variable  $X$ , the expected value of  $X$ , written  $E[X]$  is

$$E[X] = \sum_{x \in \mathcal{X}} xP(X = x)$$

Example:

$x$	$P(X)$
0	0.3
1	0.2
2	0.5

## Expected values

Given discrete random variable  $X$ , the expected value of  $X$ , written  $E[X]$  is

$$E[X] = \sum_{x \in \mathcal{X}} xP(X = x)$$

We also can talk about the expected value of functions of  $X$

$$E[f(X)] = \sum_{x \in \mathcal{X}} f(x)P(X = x)$$

## Covariance

Given two discrete r.v.'s  $X$  and  $Y$ , we define the covariance of  $X$  and  $Y$  as

$$\text{Cov}(X, Y) = E[(X - E(X))(Y - E(Y))]$$

e.g.,  $X=\text{GENDER}$ ,  $Y=\text{PLAYS\_FOOTBALL}$   
or  $X=\text{GENDER}$ ,  $Y=\text{LEFT\_HANDED}$

Remember: 
$$E[X] = \sum_{x \in \mathcal{X}} xP(X = x)$$

## Conjugate priors

- $P(\theta)$  and  $P(\theta|D)$  have the same form

**Eg. 1** Coin flip problem

Likelihood is  $\sim$  Binomial

$$P(D | \theta) = \theta^{\alpha_H} (1 - \theta)^{\alpha_T}$$

If prior is Beta distribution,

$$P(\theta) = \frac{\theta^{\beta_H-1} (1 - \theta)^{\beta_T-1}}{B(\beta_H, \beta_T)} \sim \text{Beta}(\beta_H, \beta_T)$$

Then posterior is Beta distribution

$$P(\theta|D) \sim \text{Beta}(\beta_H + \alpha_H, \beta_T + \alpha_T)$$

**For Binomial, conjugate prior is Beta distribution.**



[A. Singh]

## Conjugate priors

- $P(\theta)$  and  $P(\theta|D)$  have the same form

**Eg. 2** Dice roll problem (6 outcomes instead of 2)

Likelihood is  $\sim$  Multinomial( $\theta = \{\theta_1, \theta_2, \dots, \theta_k\}$ )

$$P(\mathcal{D} | \theta) = \theta_1^{\alpha_1} \theta_2^{\alpha_2} \dots \theta_k^{\alpha_k}$$

If prior is Dirichlet distribution,

$$P(\theta) = \frac{\prod_{i=1}^k \theta_i^{\beta_i-1}}{B(\beta_1, \dots, \beta_k)} \sim \text{Dirichlet}(\beta_1, \dots, \beta_k)$$

Then posterior is Dirichlet distribution

$$P(\theta|D) \sim \text{Dirichlet}(\beta_1 + \alpha_1, \dots, \beta_k + \alpha_k)$$

**For Multinomial, conjugate prior is Dirichlet distribution.**

[A. Singh]



## Dirichlet distribution

- number of heads in N flips of a two-sided coin
  - follows a *binomial distribution*
  - Beta is a good prior (conjugate prior for binomial)
- what if it's not two-sided, but k-sided?
  - follows a *multinomial distribution*
  - *Dirichlet* distribution is its conjugate prior

$$P(\theta_1, \theta_2, \dots, \theta_K) = \frac{1}{B(\alpha)} \prod_i^K \theta_i^{(\alpha_i-1)}$$

Lejeune Dirichlet



Johann Peter Gustav Lejeune Dirichlet

<b>Born</b>	13 February 1805 Düren, French Empire
<b>Died</b>	5 May 1859 (aged 54) Göttingen, Hanover
<b>Residence</b>	<span><span></span></span> Germany
<b>Nationality</b>	<span><span></span></span> German
<b>Fields</b>	Mathematician
<b>Institutions</b>	University of Berlin University of Breslau University of Göttingen
<b>Alma mater</b>	University of Bonn
<b>Doctoral advisor</b>	Simeon Poisson Joseph Fourier
<b>Doctoral students</b>	Ferdinand Eisenstein Leopold Kronecker Rudolf Lipschitz Carl Wilhelm Borchardt
<b>Known for</b>	Dirichlet function Dirichlet eta function